

1. **Python Basics :**
 - 1.1. **Python Basics: Google Colaboratory**
 - 1.2. **Python Basics: Basic Data Types**
 - 1.3. **Python Basics: Lists**
 - 1.4. **Python Basics: Tuples, Sets, Dictionaries**
 - 1.5. **Python Basics: Packages**
 - 1.6. **Python Basics: File Handling**
 - 1.7. **Python Basics: Classes**
 - 1.8. **Python Basics: Numpy**
 - 1.9. **Python Basics: Plotting**
2. **Expert Systems Theory :**
 - 2.1. **Expert Systems & 6 Jars: Expert Systems**
 - 2.2. **Expert Systems & 6 Jars: Say Hi To ML**
 - 2.3. **Expert Systems & 6 Jars: Introduction**
 - 2.4. **Expert Systems & 6 Jars: Data**
 - 2.5. **Expert Systems & 6 Jars: Tasks**
 - 2.6. **Expert Systems & 6 Jars: Models**
 - 2.7. **Expert Systems & 6 Jars: Loss Function**
 - 2.8. **Expert Systems & 6 Jars: Learning Algorithm**
 - 2.9. **Expert Systems & 6 Jars: Evaluation**
 - 2.10. **Expert Systems & 6 Jars: Six Jars Summary (Part 1)**
 - 2.11. **Expert Systems & 6 Jars: Six Jars Summary (Part 2)**
3. **Linear Algebra Basics :**
 - 3.1. **Vectors & Matrices: Introduction to Vectors**
 - 3.2. **Vectors & Matrices: Dot product of vectors**
 - 3.3. **Vectors & Matrices: Unit Vectors**
 - 3.4. **Vectors & Matrices: Projection of one vector onto another**
 - 3.5. **Vectors & Matrices: Angle between two vectors**
 - 3.6. **Vectors & Matrices: Why do we care about vectors?**
 - 3.7. **Vectors & Matrices: Introduction to Matrices**
 - 3.8. **Vectors & Matrices: Multiplying a vector by a matrix**
 - 3.9. **Vectors & Matrices: Multiplying a matrix by another matrix**
 - 3.10. **Vectors & Matrices: An alternate way of multiplying two matrices**
 - 3.11. **Vectors & Matrices: Why do we care about matrices?**
4. **Python Basics 2 :**
 - 4.1. **Python Basics + Linear Algebra: Google Drive and Colab Integration**
 - 4.2. **Python Basics + Linear Algebra: Pandas**
 - 4.3. **Python Basics + Linear Algebra: Python Debugger**
 - 4.4. **Python Basics + Linear Algebra: Plotting Vectors**
 - 4.5. **Python Basics + Linear Algebra: Vector Addition and Subtraction**
 - 4.6. **Python Basics + Linear Algebra: Vector Dot Product**
5. **MP Neuron :**

- 5.1. **MP Neuron: Introduction**
- 5.2. **MP Neuron: Model**
- 5.3. **MP Neuron: Data & Task**
- 5.4. **MP Neuron: Loss**
- 5.5. **MP Neuron: Learning**
- 5.6. **MP Neuron: Evaluation**
- 5.7. **MP Neuron: Geometry Basics**
- 5.8. **MP Neuron: Geometric Interpretation**
- 5.9. **MP Neuron: Summary**

- 6. **Perceptron Model :**

- 6.1. **Perceptron: Introduction**
- 6.2. **Perceptron: Data & Task**
- 6.3. **Perceptron: Model**
- 6.4. **Perceptron: Geometric Interpretation**
- 6.5. **Perceptron: Loss Function**
- 6.6. **Perceptron: Learning - General Recipe**
- 6.7. **Perceptron: Learning Algorithm**
- 6.8. **Perceptron: Learning - Why it Works?**
- 6.9. **Perceptron: Learning - Will it Always Work?**
- 6.10. **Perceptron: Evaluation**
- 6.11. **Perceptron: Summary**

- 7. **MP Neuron and Perceptron Practical :**

- 7.1. **MP Neurons & Perceptron using Python: Toy Example (Perceptron)**
- 7.2. **MP Neurons & Perceptron using Python: Loading Data**
- 7.3. **MP Neurons & Perceptron using Python: Train-Test Split**
- 7.4. **MP Neurons & Perceptron using Python: Binarization**
- 7.5. **MP Neurons & Perceptron using Python: Inference And Search**
- 7.6. **MP Neurons & Perceptron using Python: Inference**
- 7.7. **MP Neurons & Perceptron using Python: Class**
- 7.8. **MP Neurons & Perceptron using Python: Perceptron Class**
- 7.9. **MP Neurons & Perceptron using Python: Epochs**
- 7.10. **MP Neurons & Perceptron using Python: Checkpointing**
- 7.11. **MP Neurons & Perceptron using Python: Learning Rate**
- 7.12. **MP Neurons & Perceptron using Python: Weight Animation**
- 7.13. **MP Neurons & Perceptron using Python: Exercises**

- 8. **Contest 1 :**

- 8.1. **Contest - Mobile Phone Like/Dislike predictor: Introduction to Contests.**
- 8.2. **Contest - Mobile Phone Like/Dislike predictor: Creating a Kaggle account**
- 8.3. **Contest - Mobile Phone Like/Dislike predictor: Data preprocessing**
- 8.4. **Contest - Mobile Phone Like/Dislike predictor: Submitting Entries**
- 8.5. **Contest - Mobile Phone Like/Dislike predictor: Clarifications**

- 9. **Sigmoid Neuron :**

- 9.1. 6 Jars of Sigmoid Neuron (I): Recap from last lecture
 - 9.2. 6 Jars of Sigmoid Neuron (I): Revisiting limitations of perceptron model
 - 9.3. 6 Jars of Sigmoid Neuron (I): Sigmoid Model Part 1
 - 9.4. 6 Jars of Sigmoid Neuron (I): Sigmoid Model Part 2
 - 9.5. 6 Jars of Sigmoid Neuron (I): Sigmoid Model Part 3
 - 9.6. 6 Jars of Sigmoid Neuron (I): Sigmoid Model Part 4
 - 9.7. 6 Jars of Sigmoid Neuron (I): Sigmoid: Data and Tasks
 - 9.8. 6 Jars of Sigmoid Neuron (I): Sigmoid: Loss Function
 - 9.9. 6 Jars of Sigmoid Neuron (I): Learning: Intro to Learning Algorithm
 - 9.10. 6 Jars of Sigmoid Neuron (I): Learning: Learning by guessing
 - 9.11. 6 Jars of Sigmoid Neuron (I): Learning - Error Surfaces for learning
 - 9.12. 6 Jars of Sigmoid Neuron (I): Learning - Mathematical setup for the learning algorithm
 - 9.13. 6 Jars of Sigmoid Neuron (I): Learning - The math-free version of learning algorithm
 - 9.14. 6 Jars of Sigmoid Neuron (II): Learning - Introducing Taylor Series
 - 9.15. 6 Jars of Sigmoid Neuron (II): Learning - More intuitions about Taylor series
 - 9.16. 6 Jars of Sigmoid Neuron (II): Learning - Deriving the Gradient Descent Update rule
 - 9.17. 6 Jars of Sigmoid Neuron (II): Learning - The complete learning algorithm
 - 9.18. 6 Jars of Sigmoid Neuron (II): Learning - Computing Partial Derivatives
 - 9.19. 6 Jars of Sigmoid Neuron (II): Learning - Writing the code
 - 9.20. 6 Jars of Sigmoid Neuron (II): Sigmoid - Dealing with more than 2 parameters
 - 9.21. 6 Jars of Sigmoid Neuron (II): Sigmoid - Evaluation
 - 9.22. 6 Jars of Sigmoid Neuron (II): Summary and take-aways
-
- 10. Sigmoid Neuron using Python :
 - 10.1. Sigmoid Neurons Using Python (I): Plotting Sigmoid 2D
 - 10.2. Sigmoid Neurons Using Python (I): Plotting Sigmoid 3D
 - 10.3. Sigmoid Neurons Using Python (I): Plotting Loss
 - 10.4. Sigmoid Neurons Using Python (I): Contour Plot
 - 10.5. Sigmoid Neurons Using Python (I): Class
 - 10.6. Sigmoid Neurons Using Python (I): Toy Data Fit
 - 10.7. Sigmoid Neurons Using Python (I): Toy Data Plot - 1/2
 - 10.8. Sigmoid Neurons Using Python (I): Toy Data Plot - 2/2
 - 10.9. Sigmoid Neurons Using Python (II): Loading Data
 - 10.10. Sigmoid Neurons Using Python (II): Standardisation
 - 10.11. Sigmoid Neurons Using Python (II): Test/Train Split (1/2)
 - 10.12. Sigmoid Neurons Using Python (II): Test/Train Split (2/2)
 - 10.13. Sigmoid Neurons Using Python (II): Fitting Data
 - 10.14. Sigmoid Neurons Using Python (II): Loss Plot
 - 10.15. Sigmoid Neurons Using Python (II): Progress Bar
 - 10.16. Sigmoid Neurons Using Python (II): Exercises
-

11. Probability :

- 11.1. Probability Theory Introduction
- 11.2. Probability Theory: Random Variable - Intuition
- 11.3. Probability Theory: Random Variable - Formal Definition
- 11.4. Probability Theory: Random Variable - Continuous and Discrete
- 11.5. Probability Theory: Probability Distribution
- 11.6. Probability Theory: True and Predicted Distribution
- 11.7. Probability Theory: Certain Events
- 11.8. Probability Theory: Why Do we Care About Distributions

12. Information Theory & Cross Entropy :

- 12.1. Information Theory and Cross Entropy: Expectation
- 12.2. Information Theory and Cross Entropy: Information Content
- 12.3. Information Theory and Cross Entropy: Entropy
- 12.4. Information Theory and Cross Entropy: Relation To Number Of Bits
- 12.5. Information Theory and Cross Entropy: KL-Divergence and Cross Entropy
- 12.6. Information Theory and Cross Entropy: Sigmoid Neuron and Cross Entropy
- 12.7. Information Theory and Cross Entropy: Using Cross Entropy With Sigmoid Neuron
- 12.8. Information Theory and Cross Entropy: Learning Algorithm for Cross Entropy loss function
- 12.9. Information Theory and Cross Entropy: Computing partial derivatives with cross entropy loss
- 12.10. Information Theory and Cross Entropy: Code for Cross Entropy Loss function

13. Contest 2 :

- 13.1. Contest - Text, Non-Text Classification: Overview

14. Representation Power of Function :

- 14.1. Representation Power of function: Why do we need complex functions
- 14.2. Representation Power of function: Complex functions in real world examples
- 14.3. Representation Power of function: A simple recipe for building complex functions
- 14.4. Representation Power of function: Illustrative Proof of Universal Approximation Theorem
- 14.5. Representation Power of function: Summary

15. Deep Neural Networks :

- 15.1. Deep Neural Networks: Setting the context
- 15.2. Deep Neural Networks: Data and Tasks
- 15.3. Deep Neural Networks: Model: A simple deep neural network
- 15.4. Deep Neural Networks: Model: A generic deep neural network
- 15.5. Deep Neural Networks: Model: Understanding the computations in a deep neural network
- 15.6. Deep Neural Networks: Model: The output layer of a deep neural network

- 15.7. Deep Neural Networks: Model: Output layer of a multi-class classification problem
 - 15.8. Deep Neural Networks: Model: How do you choose the right network configuration
 - 15.9. Deep Neural Networks: Loss function for binary classification
 - 15.10. Deep Neural Networks: Loss function for multi-class classification
 - 15.11. Deep Neural Networks: Learning Algorithm (non-mathy version)
 - 15.12. Deep Neural Networks: Evaluation
 - 15.13. Deep Neural Networks: Summary
-
- 16. DNNs using Python :**
 - 16.1. DNNs using Python: Outline
 - 16.2. DNNs using Python: Generating Data
 - 16.3. DNNs using Python: Classification with Sigmoid Neuron
 - 16.4. DNNs using Python: Classification with FF Network
 - 16.5. DNNs using Python: Generic Class of FF Neuron
 - 16.6. DNNs using Python: Multi Class Classification with FF Network
 - 16.7. DNNs using Python: Exercise
-
- 17. Backpropagation :**
 - 17.1. **Backpropagation - the light math version: Setting the context**
 - 17.2. **Backpropagation - the light math version: Revisiting Basic Calculus**
 - 17.3. **Backpropagation - the light math version: Why do we care about the chain rule of derivatives**
 - 17.4. **Backpropagation - the light math version: Applying chain rule across multiple paths**
 - 17.5. **Backpropagation - the light math version: Applying Chain rule in a neural network**
 - 17.6. **Backpropagation - the light math version: Computing Partial Derivatives w.r.t. a weight - Part 1**
 - 17.7. **Backpropagation - the light math version: Computing Partial Derivatives w.r.t. a weight - Part 2**
 - 17.8. **Backpropagation - the light math version: Computing Partial Derivatives w.r.t. a weight - Part 3**
 - 17.9. **Backpropagation - the light math version: Computing Partial Derivatives w.r.t. a weight when there are multiple paths**
 - 17.10. **Backpropagation - the light math version: Takeaways and what next ?**
-
- 18. Backpropagation using Python :**
 - 18.1. Backpropagating using Python: Outline
 - 18.2. Backpropagating using Python: Single Weight Update
 - 18.3. Backpropagating using Python: Single Weight Training
 - 18.4. Backpropagating using Python: Multiple Weight Update
 - 18.5. **Backpropagating using Python: Visualising Outputs**
 - 18.6. **Backpropagating using Python: Visualising Weights**

- 18.7. **Backpropagating using Python: Backpropagation for Multiple Class Classification**
- 18.8. **Backpropagating using Python: Shortened Backpropagation Code**
- 18.9. **Backpropagating using Python: Exercises**

- 19. **Backpropagation 2 :**
 - 19.1. **Backpropagation - the full version: Errata from last theory slot**
 - 19.2. **Backpropagation - the full version: Setting the Context**
 - 19.3. **Backpropagation - the full version: Intuition behind backpropagation**
 - 19.4. **Backpropagation - the full version: Understanding the dimensions of gradients**
 - 19.5. **Backpropagation - the full version: Computing Derivatives w.r.t. Output Layer - Part 1**
 - 19.6. **Backpropagation - the full version: Computing Derivatives w.r.t. Output Layer - Part 2**
 - 19.7. **Backpropagation - the full version: Computing Derivatives w.r.t. Output Layer - Part 3**
 - 19.8. **Backpropagation - the full version: Quick recap of the story so far**
 - 19.9. **Backpropagation - the full version: Computing Derivatives w.r.t. Hidden Layers - Part 1**
 - 19.10. **Backpropagation - the full version: Computing Derivatives w.r.t. Hidden Layers - Part 2**
 - 19.11. **Backpropagation - the full version: Computing Derivatives w.r.t. Hidden Layers - Part 3**
 - 19.12. **Backpropagation - the full version: Computing derivatives w.r.t. one weight in any layer**
 - 19.13. **Backpropagation - the full version: Computing derivatives w.r.t. all weights in any layer**
 - 19.14. **Backpropagation - the full version: A running example of backpropagation**
 - 19.15. **Backpropagation - the full version: Summary**

- 20. **Backpropagating - the full version using Python: Outline**
 - 20.1. **Backpropagating - the full version using Python: Benefits of Vectorisation**
 - 20.2. **Backpropagating - the full version using Python: Scalar Class - Recap**
 - 20.3. **Backpropagating - the full version using Python: Vectorising weights**
 - 20.4. **Backpropagating - the full version using Python: Vectorising inputs and weights**
 - 20.5. **Backpropagating - the full version using Python: Evaluation of Classes**
 - 20.6. **Backpropagating - the full version using Python: Exercises**

- 21. **Optimisation Algorithms Theory 1:**
 - 21.1. **Variants of Gradient Descent: A quick history of DL to set the context**
 - 21.2. **Variants of Gradient Descent: Highlighting a limitation of Gradient Descent**
 - 21.3. **Variants of Gradient Descent: A deeper look into the limitation of gradient descent**
 - 21.4. **Variants of Gradient Descent: Introducing contour maps**

- 21.5. Variants of Gradient Descent: Exercise: Guess the 3D surface
- 21.6. Variants of Gradient Descent: Visualizing gradient descent on a 2D contour map
- 21.7. Variants of Gradient Descent: Intuition for momentum based gradient descent
- 21.8. Variants of Gradient Descent: Dissecting the update rule for momentum based gradient descent
- 21.9. Variants of Gradient Descent: Running and visualizing momentum based gradient descent
- 21.10. Variants of Gradient Descent: A disadvantage of momentum based gradient descent
- 21.11. Variants of Gradient Descent: Intuition behind nesterov accelerated gradient descent
- 21.12. Variants of Gradient Descent: Running and visualizing nesterov accelerated gradient descent
- 21.13. Variants of Gradient Descent: Summary and what next

22. Optimisation Algorithms Theory 2:

- 22.1. Variants of Gradient Descent: The idea of stochastic and mini-batch gradient descent
- 22.2. Variants of Gradient Descent: Running stochastic gradient descent
- 22.3. Variants of Gradient Descent: Running mini-batch gradient descent
- 22.4. Variants of Gradient Descent: Epochs and Steps
- 22.5. Variants of Gradient Descent: Why do we need an adaptive learning rate ?
- 22.6. Variants of Gradient Descent: Introducing Adagrad
- 22.7. Variants of Gradient Descent: Running and Visualizing Adagrad
- 22.8. Variants of Gradient Descent: A limitation of Adagrad
- 22.9. Variants of Gradient Descent: Running and visualizing RMSProp
- 22.10. Variants of Gradient Descent: Running and visualizing Adam
- 22.11. Variants of Gradient Descent: Summary

23. Optimisation Algorithms Practical:

- 23.1. Implementing Optimization Algorithms Using Python: Outline
- 23.2. Implementing Optimization Algorithms Using Python: Modified Sigmoid Neuron Class
- 23.3. Implementing Optimization Algorithms Using Python: Setup for Plotting
- 23.4. Implementing Optimization Algorithms Using Python: Gradient Descent Algorithm
- 23.5. Implementing Optimization Algorithms Using Python: GD Algorithm - Contour Plot
- 23.6. Implementing Optimization Algorithms Using Python: Momentum
- 23.7. Implementing Optimization Algorithms Using Python: Nesterov Accelerated GD
- 23.8. Implementing Optimization Algorithms Using Python: Mini-Batch GD
- 23.9. Implementing Optimization Algorithms Using Python: AdaGrad
- 23.10. Implementing Optimization Algorithms Using Python: RMSProp
- 23.11. Implementing Optimization Algorithms Using Python: Adam
- 23.12. Implementing Optimization Algorithms Using Python: Vectorised Class Recap
- 23.13. Implementing Optimization Algorithms Using Python: Vectorised GD Algorithms

- 23.14. Implementing Optimization Algorithms Using Python: Performance of Different Algorithms
- 23.15. Implementing Optimization Algorithms Using Python: Good solutions and Exercise

24. Activation Functions Theory:

- 24.1. Activation Functions & Initialization Methods: Setting the context
- 24.2. Activation Functions & Initialization Methods: Saturation in logistic neuron
- 24.3. Activation Functions & Initialization Methods: Zero centered functions
- 24.4. Activation Functions & Initialization Methods: Introducing Tanh and ReLU activation functions
- 24.5. Activation Functions & Initialization Methods: Tanh and ReLU Activation Functions
- 24.6. Activation Functions & Initialization Methods: Symmetry Breaking Problem
- 24.7. Activation Functions & Initialization Methods: Xavier and He initialization
- 24.8. Activation Functions & Initialization Methods: Summary and what next

25. Activation Functions Practical :

- 25.1. Hands-on - activation functions & initialization methods: Introduction and Activation Functions
- 25.2. Hands-on - activation functions & initialization methods: Activation Functions
- 25.3. Hands-on - activation functions & initialization methods: Plotting Setup
- 25.4. Hands-on - activation functions & initialization methods: Sigmoid
- 25.5. Hands-on - activation functions & initialization methods: Tanh
- 25.6. Hands-on - activation functions & initialization methods: ReLu
- 25.7. Hands-on - activation functions & initialization methods: Leaky ReLU
- 25.8. Hands-on - activation functions & initialization methods: Exercises

26. Regularization Theory:

- 26.1. Regularization: Simple v/s complex models
- 26.2. Regularization: Analysing the behavior of simple and complex models
- 26.3. Regularization: Bias and Variance
- 26.4. Regularization: Test error due to high bias and high variance
- 26.5. Regularization: Overfitting in deep neural networks
- 26.6. Regularization: A detour into hyperparameter tuning
- 26.7. Regularization: L2 regularization
- 26.8. Regularization: Dataset Augmentation and Early Stopping
- 26.9. Regularization: Summary

27. Regularization Practical :

- 27.1. Hands-on - Regularization: Outline and Libraries
 - 27.2. Hands-on - Regularization: L2 Regularisation in Code
 - 27.3. Hands-on - Regularization: Bias on Increasing Model Complexity
 - 27.4. Hands-on - Regularization: L2 Regularisation in Action
 - 27.5. Hands-on - Regularization: Adding Noise to Input Features
 - 27.6. Hands-on - Regularization: Early Stopping and Exercises
-

28. Basics of Pytorch :

- 28.1. Basics of Pytorch: Outline
 - 28.2. Basics of Pytorch: PyTorch Tensors
 - 28.3. Basics of Pytorch: Simple Tensor Operations
 - 28.4. Basics of Pytorch: NumPy vs PyTorch
 - 28.5. Basics of Pytorch: GPU PyTorch
 - 28.6. Basics of Pytorch: Automatic Differentiation
 - 28.7. Basics of Pytorch: Loss Function with AutoDiff
 - 28.8. Basics of Pytorch: Learning Loop GPU (Plays on Safari)
-

29. FNNs using Pytorch :

- 29.1. FNNs using Pytorch: Outline
 - 29.2. FNNs using Pytorch: Forward Pass With Tensors
 - 29.3. FNNs using Pytorch: Functions for Loss, Accuracy, Backpropagation
 - 29.4. FNNs using Pytorch: PyTorch Modules - NN and Optim
 - 29.5. FNNs using Pytorch: NN Sequential and Code Structure
 - 29.6. FNNs using Pytorch: GPU Execution
 - 29.7. FNNs using Pytorch: Exercises and Recap
-

30. The Convolution Operation :

- 30.1. The convolution operation: Setting the Context
 - 30.2. The convolution operation: The 1D convolution operation
 - 30.3. The convolution operation: The 2D Convolution Operation
 - 30.4. The convolution operation: Examples of 2D convolution
 - 30.5. The convolution operation: 2D convolution with a 3D filter
 - 30.6. The convolution operation: Terminology
 - 30.7. The convolution operation: Padding and Stride
-

31. Convolution to Neural Networks :

- 31.1. From convolution operation to neural networks: How is the convolution operation related to Neural Networks - Part 1
 - 31.2. From convolution operation to neural networks: How is the convolution operation related to Neural Networks - Part 2
 - 31.3. From convolution operation to neural networks: How is the convolution operation related to Neural Networks - Part 3
 - 31.4. From convolution operation to neural networks: Understanding the input/output dimensions
 - 31.5. From convolution operation to neural networks: Sparse Connectivity and Weight Sharing
 - 31.6. From convolution operation to neural networks: Max Pooling and Non-Linearities
 - 31.7. From convolution operation to neural networks: Our First Convolutional Neural Network (CNN)
 - 31.8. From convolution operation to neural networks: Training CNNs
 - 31.9. From convolution operation to neural networks: Summary and what next
-

32. CNNs in Pytorch :

- 32.1. CNNs in Pytorch: Outline
 - 32.2. CNNs in Pytorch: Loading Data Sets
 - 32.3. CNNs in Pytorch: Visualising Weights
 - 32.4. CNNs in Pytorch: Single Convolutional Layer
 - 32.5. CNNs in Pytorch: Deep CNNs
 - 32.6. CNNs in Pytorch: LeNet
 - 32.7. CNNs in Pytorch: Training Le Net
 - 32.8. CNNs in Pytorch: Visualising Intermediate Layers, Exercises
-

33. CNN Architectures :

- 33.1. CNN Architectures - Part 1: Setting the context
 - 33.2. CNN Architectures - Part 1: The Imagenet Challenge
 - 33.3. CNN Architectures - Part 1: Understanding the first layer of AlexNet
 - 33.4. CNN Architectures - Part 1: Understanding all layers of AlexNet
 - 33.5. CNN Architectures - Part 1: ZFNet
 - 33.6. CNN Architectures - Part 1: VGGNet
 - 33.7. CNN Architectures - Part 1: Summary
 - 33.8. CNN Architectures - Part 2: Setting the context
 - 33.9. CNN Architectures - Part 2: Number of computations in a convolution layer
 - 33.10. CNN Architectures - Part 2: 1x1 Convolutions
 - 33.11. CNN Architectures - Part 2: The Intuition behind GoogLeNet
 - 33.12. CNN Architectures - Part 2: The Inception Module
 - 33.13. CNN Architectures - Part 2: The GoogleNet Architecture
 - 33.14. CNN Architectures - Part 2: Average Pooling
 - 33.15. CNN Architectures - Part 2: Auxiliary Loss for training a deep network
 - 33.16. CNN Architectures - Part 2: ResNet
-

34. Building CNNs :

- 34.1. Building CNN Architectures Using Pytorch: Outline
 - 34.2. Building CNN Architectures Using Pytorch: Image Transforms
 - 34.3. Building CNN Architectures Using Pytorch: VGG
 - 34.4. Building CNN Architectures Using Pytorch: Training VGG
 - 34.5. Building CNN Architectures Using Pytorch: Pre-trained Models
 - 34.6. Building CNN Architectures Using Pytorch: Checkpointing Models
 - 34.7. Building CNN Architectures Using Pytorch: ResNet
 - 34.8. Building CNN Architectures Using Pytorch: Inception Part 1
 - 34.9. Building CNN Architectures Using Pytorch: Inception Part 2
 - 34.10. Building CNN Architectures Using Pytorch: Exercises
-

35. Visualising CNNs :

- 35.1. Visualising CNNs: Receptive field of a neuron
- 35.2. Visualising CNNs: Identifying images which cause certain neurons to fire
- 35.3. Visualising CNNs: Visualising filters
- 35.4. Visualising CNNs: Occlusion experiments
- 35.5. Visualising CNNs Using Python: Outline

- 35.6. Visualising CNNs Using Python: Custom Torchvision Dataset
- 35.7. Visualising CNNs Using Python: Visualising inputs
- 35.8. Visualising CNNs Using Python: Occlusion
- 35.9. Visualising CNNs Using Python: Visualising filters
- 35.10. Visualising CNNs Using Python: Visualising filters - code

36. Batch Normalization and Dropout :

- 36.1. Batch Normalization and Dropout: Normalizing inputs
- 36.2. Batch Normalization and Dropout: Why should we normalize the inputs
- 36.3. Batch Normalization and Dropout: Batch Normalization
- 36.4. Batch Normalization and Dropout: Learning Mu and Sigma
- 36.5. Batch Normalization and Dropout: Ensemble Methods
- 36.6. Batch Normalization and Dropout: The idea of dropout
- 36.7. Batch Normalization and Dropout: Training without dropout
- 36.8. Batch Normalization and Dropout: How does weight sharing help ?
- 36.9. Batch Normalization and Dropout: Using dropout at test time
- 36.10. Batch Normalization and Dropout: How does dropout act as a regularizer ?
- 36.11. Batch Normalization and Dropout: Summary and what next ?

37. Batch Normalization and Dropout using Python :

- 37.1. Batch Normalization and Dropout Using Python: Outline and Dataset
- 37.2. Batch Normalization and Dropout Using Python: Batch Norm Layer
- 37.3. Batch Normalization and Dropout Using Python: Batch Norm Visualisation
- 37.4. Batch Normalization and Dropout Using Python: Batch Norm 2d
- 37.5. Batch Normalization and Dropout Using Python: Dropout layer
- 37.6. Batch Normalization and Dropout Using Python: Dropout Visualisation and Exercises

38. Hyperparameter Tuning and MLFlow:

- 38.1. Hyperparameter Tuning and MLFlow: Outline
- 38.2. Hyperparameter Tuning and MLFlow: Colab on Local Runtime
- 38.3. Hyperparameter Tuning and MLFlow: MLFlow installation and basic usage
- 38.4. Hyperparameter Tuning and MLFlow: Hyperparameter Tuning
- 38.5. Hyperparameter Tuning and MLFlow: Refined Search for Hyperparameters
- 38.6. Hyperparameter Tuning and MLFlow: Logging Image Artifacts
- 38.7. Hyperparameter Tuning and MLFlow: Logging and Loading Models
- 38.8. Hyperparameter Tuning and MLFlow: One Last Visualisation

39. Sequence Learning Problems :

- 39.1. Sequence Learning Problems: Setting the context
- 39.2. Sequence Learning Problems: Introduction to sequence learning problems
- 39.3. Sequence Learning Problems: Some more examples of sequence learning problems
- 39.4. Sequence Learning Problems: Sequence learning problems using video and speech data

39.5.	Sequence Learning Problems: A wishlist for modelling sequence learning problems
39.6.	Sequence Learning Problems: Intuition behind RNNs - Part 1
39.7.	Sequence Learning Problems: Intuition behind RNNs - Part 2
39.8.	Sequence Learning Problems: Introducing RNNs
39.9.	Sequence Learning Problems: Summary and what next
40.	Recurrent Neural Networks :
40.1.	Recurrent Neural Networks: Setting the context
40.2.	Recurrent Neural Networks: Data and Tasks - Sequence Classification - Part 1
40.3.	Recurrent Neural Networks: Data and Tasks - Sequence Classification - Part 2
40.4.	Recurrent Neural Networks: A clarification about padding
40.5.	Recurrent Neural Networks: Data and Tasks - Sequence Labelling
40.6.	Recurrent Neural Networks: Model
40.7.	Recurrent Neural Networks: Loss Function
40.8.	Recurrent Neural Networks: Learning Algorithm
40.9.	Recurrent Neural Networks: Learning Algorithm - Derivatives w.r.t. V
40.10.	Recurrent Neural Networks: Learning Algorithm - Derivatives w.r.t. W
40.11.	Recurrent Neural Networks: Evaluation
40.12.	Recurrent Neural Networks: Summary and what next
41.	Vanishing and Exploding Gradients :
41.1.	Vanishing and Exploding Gradients: Revisiting the gradient wrt W
41.2.	Vanishing and Exploding Gradients: Zooming into one element of the chain rule - Part 1
41.3.	Vanishing and Exploding Gradients: Zooming into one element of the chain rule - Part 2
41.4.	Vanishing and Exploding Gradients: A small detour to calculus
41.5.	Vanishing and Exploding Gradients: Looking at the magnitude of the derivative
41.6.	Vanishing and Exploding Gradients: Exploding and vanishing gradients
41.7.	Vanishing and Exploding Gradients: Summary and what next
42.	LSTMs and GRUs:
42.1.	LSTMs and GRUs: Dealing with longer sequences
42.2.	LSTMs and GRUs: The white board analogy
42.3.	LSTMs and GRUs: Real world example of longer sequences
42.4.	LSTMs and GRUs: Going back to RNNs
42.5.	LSTMs and GRUs: Selective Write - Part 1
42.6.	LSTMs and GRUs: Selective Write - Part 2
42.7.	LSTMs and GRUs: Selective Read
42.8.	LSTMs and GRUs: Selective forget
42.9.	LSTMs and GRUs: An example computation with LSTMs
42.10.	LSTMs and GRUs: Gated recurrent units
42.11.	LSTMs and GRUs: Summary and what next
43.	Sequence Models in Pytorch :

- 43.1. Sequence Models in PyTorch: Outline
 - 43.2. Sequence Models in PyTorch: Dataset and Task
 - 43.3. Sequence Models in PyTorch: RNN Model
 - 43.4. Sequence Models in PyTorch: Inference on RNN
 - 43.5. Sequence Models in PyTorch: Training RNN
 - 43.6. Sequence Models in PyTorch: Training Setup
 - 43.7. Sequence Models in PyTorch: LSTM
 - 43.8. Sequence Models in PyTorch: GRU and Exercises
-
- 44. Addressing the problem of vanishing and exploding gradients :**
- 44.1. Addressing the problem of vanishing and exploding gradients: Quick Recap
 - 44.2. Addressing the problem of vanishing and exploding gradients: Intuition: How gates help to solve the problem of vanishing gradients
 - 44.3. Addressing the problem of vanishing and exploding gradients: Revisiting vanishing gradients in RNNs
 - 44.4. Addressing the problem of vanishing and exploding gradients: Dependency diagram for LSTMs
 - 44.5. Addressing the problem of vanishing and exploding gradients: Computing the gradient
 - 44.6. Addressing the problem of vanishing and exploding gradients: When do the gradients vanish?
 - 44.7. Addressing the problem of vanishing and exploding gradients: Dealing with exploding gradients
 - 44.8. Addressing the problem of vanishing and exploding gradients: Summary and what next
-
- 45. Batching for Sequence Models in Pytorch :**
- 45.1. Batching for Sequence Models in Pytorch: Overview
 - 45.2. Batching for Sequence Models in Pytorch: Recap on Sequence Models
 - 45.3. Batching for Sequence Models in Pytorch: Batching for Sequence Models
 - 45.4. Batching for Sequence Models in Pytorch: Padding Vector Representations
 - 45.5. Batching for Sequence Models in Pytorch: Packing in PyTorch
 - 45.6. Batching for Sequence Models in Pytorch: Training with Batched Input
-
- 46. Neural Encoders and Decoders :**
- 46.1. Neural Encoders and Decoders: Setting the context
 - 46.2. Neural Encoders and Decoders: Revisiting the task of language modelling
 - 46.3. Neural Encoders and Decoders: Using RNNs for language modelling
 - 46.4. Neural Encoders and Decoders: Introducing Encoder Decoder Model
 - 46.5. Neural Encoders and Decoders: Connecting encoder decoder models to the six jars
 - 46.6. Neural Encoders and Decoders: A compact notation for RNNs, LSTMs and GRUs
 - 46.7. Neural Encoders and Decoders: Encoder decoder model for image captioning
 - 46.8. Neural Encoders and Decoders: Six jars for image captioning

46.9.	Neural Encoders and Decoders: Encoder decoder for Machine translation
46.10.	Neural Encoders and Decoders: Encoder decoder model for transliteration
46.11.	Neural Encoders and Decoders: Summary
47.	Attention Mechanism :
47.1.	Attention Mechanism: Motivation for attention mechanism
47.2.	Attention Mechanism: Attention mechanism with an oracle
47.3.	Attention Mechanism: A model for attention
47.4.	Attention Mechanism: The attention function
47.5.	Attention Mechanism: Machine translation with attention
47.6.	Attention Mechanism: Summary and what next
48.	Encoder Decoder Models and Attention Mechanism using PyTorch :
48.1.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Outline
48.2.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Data set and Task
48.3.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Data Ingestion - XML processing
48.4.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Encoder Decoder Model - 1
48.5.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Encoder Decoder Model - 2
48.6.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Adding Attention - 1
48.7.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Adding Attention - 2
48.8.	Encoder Decoder Models and Attention Mechanism Using Pytorch: Model Evaluation and Exercises
49.	RCNN :
49.1.	RCNN: More clarity on regression
49.2.	RCNN: Setting the context
49.3.	RCNN: A typical pipeline for object detection
49.4.	RCNN: Region Proposal
49.5.	RCNN: Feature Extraction
49.6.	RCNN: Classification
49.7.	RCNN: Regression
49.8.	RCNN: Training
50.	YOLO :
50.1.	YOLO: Introduction to YOLO
50.2.	YOLO: The Output of YOLO
50.3.	YOLO: Training
50.4.	Object Detection: Summary and what next
51.	Capstone :
51.1.	Capstone Project: Project Details

Watch lectures and then try them out in separate lab(sort of) time. Organise your schedule according to that. Because most of the theory part we already know, it's the implementation part we need to sort and that can be done with better resonance of time meant for research and implementation. There can be different pace for the two imo. The lectures can be seen as their input to you and then your experiments should be independent of the exact lecture words. Rather they should be more open ended and broad.

Spend good time on RL as well. Learning that and being able to implement those concepts could really help you moving forward. Other concepts can be picked up with time, on weekends or at times when you want to do something absolutely different. You have a lot of options to pick out from for that. Flying car nanodegree is the cornerstone of all that and we are finally finding something that has been scalable, is scalable and will be scalable in the foreseeable future, unlike RL or DL and so on, and that is robotics. Spend more time on robotics in the coming future and that will be very rewarding imo. Work on NLP and CV as well.

Completing deep learning should be a good priority, because once that is done, then we can go about miscellaneous sources like research papers and then explore them for other stuff that is blocked right now for some reason. Like computer vision and natural language processing is something we should properly focus on working and exploring and then implementing stuff.

Let's split the course into theory and practical. Theory includes all lectures and their note making and then there is the practical slot wherein all related exploration and basic code implementation needs to be done related to the code and the concepts. I really think that having a mix in the practical part would not be as such a bad thing, and if we can finish with the theory part by the time the btp is done with, I think we would have a good amount of time before the year is up to finish the practical aspect of exploration and learning things related to deep learning. And then from next semester we can properly focus on the application part of things.

We need to be okay with the process being slow and then going about getting better every day rather than wanting fast progress. We need to manage time better, and then make time for both the theory and practical parts of the course, along with other lecture series and resources. Reading, exploring and coding stuff is the only way to get better at this stuff, and to breaking inhibitions.

We need to be okay investing some good amount of time for the same and not rush through stuff. Rather it should be a mix of new stuff and then exploring on the knowledge already compiled, so as to grow in different forms, and not limiting ourselves in any way. But this process requires time and patience to pay off for sure and also some long nights along with perfect discipline.

This should be a mix of reading, watching independent videos and lecture series as well. But one thing we need to be sure is to complete whatever we pick up. The big problem with me is that I pick up stuff and then I don't complete it properly and then I get bored and pick something else up as well. Which really hurts the entire process. Rather having a mix of things and then focusing them should be a good way to go at things.

Thinking that you are not doing adequate work is also one reason for unwarranted stress that can surely be avoided. Focus on 1-2 activities and finish them properly and then pick up something new. But in the 1-2 activities make sure that you immerse yourself properly. And also remember that you can't consume everything, there will always be something that you can't get to and something that you will miss for sure. And also there are things which you would benefit more from by focussing on now, rather than something else.

-
1. Machine Learning
 - a. Coding Ninjas
 2. Deep Learning
 - a. One Fourth Labs Course
-

- b. Lecture Series by fast.ai
- 3. Self Driving Cars
 - a. Course by Coursera
 - b. ARES pipeline research and development
- 4. Reinforcement Learning
 - a. Udacity Course
 - b. CS285 by UC Berkeley
- 5. Robotics
- 6. Aerial Vehicles
- 7. Random Algorithms
- 8. Mathematics
- 9. Natural Language Processing
- 10. Computer Vision
- 11. Web Development
 - a. Colt Steele Course
- 12. Open Source Development
- 13. Google Summer of Code

Deep Learning, do properly and Web Development do properly. Extending them won't be that hard. Mix it up in between. Like on weekends, do something something entirely different.

Implementing papers, reading one paper per week and making notes on them.